

# Static Analysis Debugging with Symbolic Execution

**Theodoros Kasampalis, Sandeep Dasgupta**

11th August 2015

# Outline

- 1 Static Analysis
- 2 Debugging a Static Analysis Implementation
- 3 Related Work
- 4 Background
- 5 Our Idea
- 6 System Status Overview
- 7 Implementation
- 8 Questions?

# Static Analysis

- Infer source code properties without execution

# Static Analysis

- Infer source code properties without execution
- Examples:

# Static Analysis

- Infer source code properties without execution
- Examples:
  - Pointer Analysis

# Static Analysis

- Infer source code properties without execution
- Examples:
  - Pointer Analysis
  - Liveness Analysis

# Static Analysis

- Infer source code properties without execution
- Examples:
  - Pointer Analysis
  - Liveness Analysis
- Applications:

# Static Analysis

- Infer source code properties without execution
- Examples:
  - Pointer Analysis
  - Liveness Analysis
- Applications:
  - Compilers



# Static Analysis

- Infer source code properties without execution
- Examples:
  - Pointer Analysis
  - Liveness Analysis
- Applications:
  - Compilers
  - Security

# Static Analysis

- Infer source code properties without execution
- Examples:
  - Pointer Analysis
  - Liveness Analysis
- Applications:
  - Compilers
  - Security
  - Software Engineering

# Static Analysis

- Infer source code properties without execution
- Examples:
  - Pointer Analysis
  - Liveness Analysis
- Applications:
  - Compilers
  - Security
  - Software Engineering
- Inferred properties true for any execution

# Outline

- 1 Static Analysis
- 2 Debugging a Static Analysis Implementation**
- 3 Related Work
- 4 Background
- 5 Our Idea
- 6 System Status Overview
- 7 Implementation
- 8 Questions?

# Debugging a Static Analysis Implementation

- Semantic bugs

# Debugging a Static Analysis Implementation

- Semantic bugs
  - no crash

# Debugging a Static Analysis Implementation

- Semantic bugs
  - no crash
  - erroneous results

# Debugging a Static Analysis Implementation

- Semantic bugs
  - no crash
  - erroneous results
- Effect visible in client code



# Debugging a Static Analysis Implementation

- Semantic bugs
  - no crash
  - erroneous results
- Effect visible in client code
  - Hard to trace back

# Debugging a Static Analysis Implementation

- Semantic bugs
  - no crash
  - erroneous results
- Effect visible in client code
  - Hard to trace back
  - Not reliable

# Debugging a Static Analysis Implementation

- Semantic bugs
  - no crash
  - erroneous results
- Effect visible in client code
  - Hard to trace back
  - Not reliable
- Static analysis specific tests

# Debugging a Static Analysis Implementation

- Semantic bugs
  - no crash
  - erroneous results
- Effect visible in client code
  - Hard to trace back
  - Not reliable
- Static analysis specific tests
  - small regression tests

# Outline

- 1 Static Analysis
- 2 Debugging a Static Analysis Implementation
- 3 Related Work**
- 4 Background
- 5 Our Idea
- 6 System Status Overview
- 7 Implementation
- 8 Questions?

# Related Work

- Compiler testing through miscompilation detection:

# Related Work

- Compiler testing through miscompilation detection:
  - Test suites (LLVM test suite)

# Related Work

- Compiler testing through miscompilation detection:
  - Test suites (LLVM test suite)
  - Randomly generated tests (Csmith, PLDI 11)



# Related Work

- Compiler testing through miscompilation detection:
  - Test suites (LLVM test suite)
  - Randomly generated tests (Csmith, PLDI 11)
  - Equivalence Modulo Inputs (Orion, PLDI 14)

# Related Work

- Compiler testing through miscompilation detection:
  - Test suites (LLVM test suite)
  - Randomly generated tests (Csmith, PLDI 11)
  - Equivalence Modulo Inputs (Orion, PLDI 14)
- Dynamic alias analysis error detection (NeonGoby, FSE 13)

# Related Work

- Compiler testing through miscompilation detection:
  - Test suites (LLVM test suite)
  - Randomly generated tests (Csmith, PLDI 11)
  - Equivalence Modulo Inputs (Orion, PLDI 14)
- Dynamic alias analysis error detection (NeonGoby, FSE 13)
- Symbolic execution (KLEE, OSDI 08)

# Related Work

- Compiler testing through miscompilation detection:
  - Test suites (LLVM test suite)
  - Randomly generated tests (Csmith, PLDI 11)
  - Equivalence Modulo Inputs (Orion, PLDI 14)
- Dynamic alias analysis error detection (NeonGoby, FSE 13)
- Symbolic execution (KLEE, OSDI 08)
- Concolic execution (zesti, ICSE 12 - SAGE, ICSE13)

# Outline

- 1 Static Analysis
- 2 Debugging a Static Analysis Implementation
- 3 Related Work
- 4 Background**
- 5 Our Idea
- 6 System Status Overview
- 7 Implementation
- 8 Questions?

# Array Out of Bounds Bug

```
1. int v[100];

2. void f(int x) {
3.     if (x > 99)
4.         x = 99;

5.     v[x] = 0;
6. }

7. int main(int argc, char **argv) {
8.     int x = 50;
9.     f(x);

10.    return 0;
11. }
```

# Symbolic Execution with KLEE

```
1. int v[100];  
  
2. void f(int x) {  
3.     if (x > 99)  
4.         x = 99;  
  
5.     v[x] = 0;  
6. }  
  
7. int main(int argc, char **argv) {  
8.     int x;  
9.     klee_make_symbolic(&x, sizeof(x), "X");  
10.    f(x);  
  
11.    return 0;  
12. }
```

# Concolic Execution with zesti

```
1. int v[100];  
  
2. void f(int x) {  
3.     if (x > 99)  
4.         x = 99;  
  
5.     v[x] = 0;  
6. }  
  
7. int main(int argc, char **argv) {  
8.     int x = 50;  
9.     klee_make_symbolic(&x, sizeof(x), "X");  
10.    f(x);  
  
11.    return 0;  
12. }
```



# Concolic Execution with zesti

```
1. int v[100];  
  
2. void f(int x) {  
3.     if (x > 99)  
4.         x = 99;  
  
5.     v[x] = 0;  
6. }  
  
7. int main(int argc, char **argv) {  
8.     int x = 100;  
9.     klee_make_symbolic(&x, sizeof(x), "X");  
10.    f(x);  
  
11.    return 0;  
12. }
```

# Outline

- 1 Static Analysis
- 2 Debugging a Static Analysis Implementation
- 3 Related Work
- 4 Background
- 5 Our Idea**
- 6 System Status Overview
- 7 Implementation
- 8 Questions?

# Our Idea

- Check inferred properties during symbolic execution

# Our Idea

- Check inferred properties during symbolic execution
  - Apply analysis to an input program

# Our Idea

- Check inferred properties during symbolic execution
  - Apply analysis to an input program
  - Symbolically execute the input program

# Our Idea

- Check inferred properties during symbolic execution
  - Apply analysis to an input program
  - Symbolically execute the input program
  - Check whether inferred properties hold

# Our Idea

- Check inferred properties during symbolic execution
  - Apply analysis to an input program
  - Symbolically execute the input program
  - Check whether inferred properties hold
  
- Direct testing of static analysis code

# Our Idea

- Check inferred properties during symbolic execution
  - Apply analysis to an input program
  - Symbolically execute the input program
  - Check whether inferred properties hold
- Direct testing of static analysis code
- Static analysis inferences checked thoroughly



# Our Idea

- Check inferred properties during symbolic execution
  - Apply analysis to an input program
  - Symbolically execute the input program
  - Check whether inferred properties hold
- Direct testing of static analysis code
- Static analysis inferences checked thoroughly
  - High path coverage of the input program

# Our Idea

- Check inferred properties during symbolic execution
  - Apply analysis to an input program
  - Symbolically execute the input program
  - Check whether inferred properties hold
- Direct testing of static analysis code
- Static analysis inferences checked thoroughly
  - High path coverage of the input program
  - Big input program size

# Outline

- 1 Static Analysis
- 2 Debugging a Static Analysis Implementation
- 3 Related Work
- 4 Background
- 5 Our Idea
- 6 System Status Overview**
- 7 Implementation
- 8 Questions?

# System Status Overview

- Implementation for checking an LLVM Alias Analysis (including tbaa, basicaa)

# System Status Overview

- Implementation for checking an LLVM Alias Analysis (including tbaa, basicaa)
- Checks incorporated within zesti

# System Status Overview

- Implementation for checking an LLVM Alias Analysis (including tbaa, basicaa)
- Checks incorporated within zesti
- Checks on all loads

# System Status Overview

- Implementation for checking an LLVM Alias Analysis (including tbaa, basicaa)
- Checks incorporated within zesti
- Checks on all loads
- Pointer dereferences marked sensitive

# System Status Overview

- Implementation for checking an LLVM Alias Analysis (including tbaa, basicaa)
- Checks incorporated within zesti
- Checks on all loads
- Pointer dereferences marked sensitive
- Reachability analysis for inputs affecting pointer values [under implementation]



# System Status Overview

- Implementation for checking an LLVM Alias Analysis (including tbaa, basicaa)
- Checks incorporated within zesti
- Checks on all loads
- Pointer dereferences marked sensitive
- Reachability analysis for inputs affecting pointer values [under implementation]
- Testing with LLVM test suite programs

# Outline

- 1 Static Analysis
- 2 Debugging a Static Analysis Implementation
- 3 Related Work
- 4 Background
- 5 Our Idea
- 6 System Status Overview
- 7 Implementation**
- 8 Questions?

# Symbolic Execution

- Symbolic execution using klee
- Migration from Klee to Zesti (a variant of klee)

# Debugger Logic for Pointer Analysis

```
foreach(loadI = load instructions) {
  base_address = 'base address' of the loadI
  foreach( 'pointer' in the same function scope as the load instruction) {
    result = mustAlias_OR_mayNOTAlias('base_address', 'pointer') // Querying the alias analysis.
    if( result == must-alias) {
      if ('base_pointer' and 'pointer' DO NOT point to the same run-time memory object) {
        error
      }
    }
    if (result == mayNot-alias) {
      if ('base_pointer' and 'pointer' point to the same run-time memory object) {
        error
      }
    }
  }
}
```

# Implicitly adding klee\_assumes

```
struct S {
    int member;
};
struct S data[] =
{
    { 1,2 },
    { 3,4 },
};
int main(int argc, char** argv) {

    int x= 0 ;
    struct S* z;

    klee_make_symbolic(&x, sizeof(x), "X");
    /*
    ** Without the following klee_assume, the dereference z->x gets resolved to many
    ** spurious memory objects.
    ** Generated in-bound constraints on the fly to prevent this.
    */
    klee_assume(x >= 0 & x <= 1 );

    z = &data[x];
    ... = z->member ;

    return 0;
}
```

# Importance of choosing a variable as symbolic

```
1. int main() {
2.     int x=1 , y=2;
3.     int* p = (int *)malloc(sizeof(int));

4.     klee_make_symbolic(&x, sizeof(x), "x");
5.     klee_make_symbolic(&y, sizeof(y), "y");
   /*
   ** If we skip to make y symbolic, then we may miss the
   ** opportunity of catching a potential pointer analysis
   ** bug. For ex. what if the pointer analysis infers that
   ** *p and the heap object at line 7 mayNOT alias.
   */

   if(0 != x*y) {
6.       p = (int *)malloc(4);
   } else {
7.       if(y == 0) {
           p = (int *)malloc(4);
       }
   }
8.     return *p;
}
```

# Which variables to make symbolic

- Explicitly specifying which variables to make symbolic is difficult.
  - Instrumented the code by inserting appropriate `klee_make_symbolic`.
  - Reachability Analysis to figure out candidates to be made symbolic.

# Bugs Found

```
/* The bug shows up when there is a must alias check between
** x (at line 1) and the bitcast of x (at line 3).
*/
int main(int argc, char **argv) {
    int *A[5];
    for (int i = 0; i < 5; ++i) {
        A[i] = (int*) malloc((i+1)*sizeof(int));
    }

    int *x, a;
    char *y;

    for (int i = 0; i < 5; ++i) {
1.   x = A[i];
2.   a = *x;
3.   y = (char *) x;
    }
    return *y;
}
```



# Outline

- 1 Static Analysis
- 2 Debugging a Static Analysis Implementation
- 3 Related Work
- 4 Background
- 5 Our Idea
- 6 System Status Overview
- 7 Implementation
- 8 Questions?**

